Task 6 Output Comparison:

Before Fine-Tuning: The base model outputs generic text or repeats lines without summarizing.

After Fine-Tuning: The LoRA model generates accurate, concise and succinct summaries capturing the conversation's factual meaning

Task 7 Rank Experiment r=4 vs r=16:

Parameter Capacity: Increasing the rank from r=4 to r=16 quadruples the number of trainable parameters from 172k to 688k. This gives the model a larger capacity to store and adapt to patterns in the dataset.

Output Quality: Because of the higher rank, the r=16 model captures subtle conversational nuances better, leading to improved consistency, less repetition.